

Classification- Visual Recognition

1. Introduction
2. Supervised learning
3. Perceptrons
- 4. SVM classifiers**
5. Datasets and evaluation

SVM

Notations:

- Image/Patterns $\mathbf{x} \in \mathbf{X}$
- Φ : function transforming the patterns into feature vectors $\Phi(x)$
- $\langle \cdot, \cdot \rangle$ dot product in the feature space endowed by $\Phi(\cdot)$
- Classes $y = \pm 1$

Early kernel classifiers derived from the perceptron [Rosenblatt58]:

- taking the sign of a linear discriminant function:

$$f(\mathbf{x}) = \langle \mathbf{w}, \Phi(\mathbf{x}) \rangle + b$$

- Classifiers called Φ -machines



SVM

- Question: how to find/estimate f ?
 - Feature function Φ usually hand-chosen for each problem
 - Several Φ for image processing like BoW
 - w and b : parameters to be determined


$$f(x) = \langle w, \Phi(x) \rangle + b$$

- Learning algorithm on a set of training examples:
 $\mathcal{A} = (x_1, y_1) \cdots (x_n, y_n)$

Which hyperplane ? w ? b ?



SVM



SVM optimization: maximizing the margin between + and -

Def.: Margin = distance between the hyperplanes $f(x) = 1$ and $f(x) = -1$ (dashed lines in Figure).

Intuitively, a classifier with a larger margin is more robust to fluctuations



Hard Margin

Final expression for the Hard Margin SVM optimization:

$$\min_{w,b} P(w,b) = \frac{1}{2} \|w\|^2 \quad \text{with} \quad \forall i \quad y_i f(x_i) \geq 1$$

SVM

- Hard Margin: OK if data are linearly separated
- Otherwise: noisy data (in red) disrupt the optim.
- Solution: Soft SVM



SVM: Soft Margin

Introducing the slack variables ξ_i , one usually gets rid of the inconvenient max of the loss and rewrite the problem as

$$\min_{w,b} P(w,b) = \frac{1}{2} \|w\|^2 + C \sum_{i=1}^n \xi_i \quad \text{with} \quad \begin{cases} \forall i & y_i f(x_i) \geq 1 - \xi_i \\ \forall i & \xi_i \geq 0 \end{cases}$$


For very large values of the hyper-parameter C, **Hard Margin** case:

- Minimization of $\|w\|$ (ie margin maximization) under the constraint that all training examples are correctly classified with a loss equal to zero.

Smaller values of C relax this constraint: **Soft Margin** case

- SVMs that produces markedly better results on noisy problems.

SVM learning scheme



Equivalently, minimizing the following objective function in feature space with the hinge loss function:

$$\ell(y_i f(x_i)) = \max(0, 1 - y_i f(x_i))$$

$$\min_{w,b} P(w,b) = \frac{1}{2} \|w\|^2 + C \sum_{i=1}^n \ell(y_i f(x_i))$$

Regularization

Margin
Maximization

Data
fitting

Constraint
satisfaction

Learning SVMs: Primal/Dual

- In practice: Convex optimization problem
 - Primal optimization: $f(x) = \langle w, \Phi(x) \rangle + b$
 - Dual optimization: learning SVMs can be achieved by solving the dual of this convex optimization problem
- Dual (using Lagrangian and $k(x_i, x_j) = \langle \Phi(x_i), \Phi(x_j) \rangle$):
 - For Hard Margin:

$$\max_{\alpha} \mathcal{L}(\alpha) = \sum_{i=1} \alpha_i - \frac{1}{2} \sum_{i,j} y_i y_j \alpha_i \alpha_j k(x_i, x_j) \quad \text{with} \quad \begin{cases} \sum_i \alpha_i y_i = 0 \\ 0 \leq \alpha_i \end{cases}$$

- For Soft Margin:

$$\max_{\alpha} \mathcal{L}(\alpha) = \sum_{i=1} \alpha_i - \frac{1}{2} \sum_{i,j} y_i y_j \alpha_i \alpha_j k(x_i, x_j) \quad \text{with} \quad \begin{cases} \sum_i \alpha_i y_i = 0 \\ 0 \leq \alpha_i \leq C \end{cases}$$

SVM optimization

Standard equivalent formulation without enforcing α_i to be positive:

- Optimization on coefficients α_i of the SVM kernel expansion $f(x) = \sum_{i=1}^n \alpha_i k(x, x_i) + b$ by defining the dual objective function:

$$D(\boldsymbol{\alpha}) = \sum_i \alpha_i y_i - \frac{1}{2} \sum_{i,j} \alpha_i \alpha_j k(x_i, x_j)$$

- and solving the SVM dual *Quadratic Programming* (QP) problem.

$$\max_{\boldsymbol{\alpha}} D(\boldsymbol{\alpha}) \quad \text{with} \quad \left\{ \begin{array}{l} \sum_i \alpha_i = 0 \\ A_i \leq \alpha_i \leq B_i \\ A_i = \min(0, Cy_i) \\ B_i = \max(0, Cy_i) \end{array} \right.$$



Classification pipeline

To summarize on SVM :

Solving equation: SVM

Support Vector Machines (SVM) defined by three incremental steps:

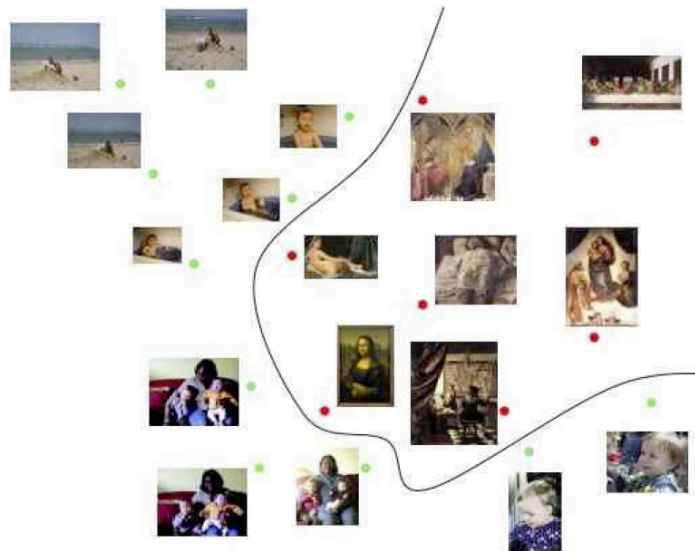
1. [Vapnik63]: linear classifier / separates the training examples with the **widest margin** => Optimal Hyperplane



Solving equation: SVM

Support Vector Machines (SVM) defined by three incremental steps:


1. [Vapnik63]: linear classifier / separates the training examples with the widest margin =>Optimal Hyperplane
2. [Guyon93] Optimal Hyperplane built in the feature space induced by a kernel function



Solving equation: SVM

Support Vector Machines (SVM) defined by three incremental steps:

1. [Vapnik63]: linear classifier / separates the training examples with the widest margin =>Optimal Hyperplane
2. [Guyon93] Optimal Hyperplane built in the feature space induced by a kernel function
3. [Cortes95] soft version: noisy problems addressed by allowing some examples to violate the margin constraint



Appendix: Solving SVM

- Min P or Max D
 - => QP (Quadratic programming) family optimization
- Good news: efficient batch numerical algorithms have been developed to solve the specific SVM QP problem (hinge loss, convex objective,...)
- Some strategies (exploiting specif.):
 - Conjugate Gradient method [Vapnik]
 - Sequential Minimal Optimization (SMO) [platt].
- In both methods successive searches along well chosen directions
- Some famous SVM solvers like SVMLight [Joachims] or SVMTorch propose to use *decomposition* algorithms to define such directions
- SVMstruct (for structured outputs)
- State-of-the-art implementation of SMO: [libsvm] => used in tutorials
- LibLinear bib for primal optim (with MATLAB)


SMO algo for SVM optimization

1. Set $\alpha \leftarrow \mathbf{0}$ and compute the initial gradient \mathbf{g} of $D(\alpha)$
2. Choose a τ -violating pair(*) (i, j) Stop if no such pair exists
3. $\lambda \leftarrow \min \left\{ \frac{g_i - g_j}{k_{ii} + k_{jj} - 2k_{ij}}, B_i - \alpha_i, \alpha_j - A_j \right\}$
4. $\alpha_i \leftarrow \alpha_i + \lambda$, $\alpha_j \leftarrow \alpha_j - \lambda$
5. $g_s \leftarrow g_s - \lambda(k_{is} - k_{js}) \quad \forall s \in \{1 \dots n\}$
6. Return to step 2

(*) pairs in +1/-1 with significant diff of gradients

A ways to easily satisfy the null sum coeff constraint

Classification pipeline



Classification- Visual Recognition

1. Introduction
2. Supervised learning
3. Perceptrons
4. SVM classifiers
5. **Datasets and evaluation**